

Problem Statement 1: - Mandatory

Explore the given data set with EDA techniques and build a suitable model for predicting whether the salary of the person is >50k or not and visualize the results.

NOTE: the algorithm used for the model must be built from the scratch.

PROBLEM STATEMENT AND ANALYSIS:

Logistic regression is the appropriate regression analysis to conduct when the dependent variable is dichotomous (binary). Like all regression analyses, the logistic regression is a predictive analysis. Logistic regression is used to describe data and to explain the relationship between one dependent binary variable and one or more nominal, ordinal, interval or ratio-level independent variables.

A support vector machine (SVM) is machine learning algorithm that analyzes data for classification and regression analysis. SVM is a supervised learning method that looks at data and sorts it into one of two categories. An SVM outputs a map of the sorted data with the margins between the two as far apart as possible. SVMs are used in text categorization, image classification, handwriting recognition and in the sciences. A support vector machine is also known as a support vector network (SVN).

Random Forest is a popular machine learning algorithm that belongs to the supervised learning technique. "Random Forest is a classifier that contains a number of decision trees on various subsets of the given dataset and takes the average to improve the predictive accuracy of that dataset." Instead of relying on one decision tree, the random forest takes the prediction from each tree and based on the majority votes of predictions, and it predicts the final output. The greater number of trees in the forest leads to higher accuracy and prevents the problem of overfitting.

CODE:

```
#DATA PREPROCESSING
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing
dt = pd.read_csv('train.csv')
dt.head()
dt.info()
dt.shape
dt.describe()
```

```
dt.isnull().sum()
```

```
train=dt.dropna()  
train.isnull().sum()
```

```
sns.heatmap(train.corr(),cmap='coolwarm')
```

```
#encoding
```

```
train.columns = ['age', 'workclass', 'fnlwgt', 'education', 'education.num', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'capiat.gain', 'capital.loss', 'hours.per.week', 'native.country', 'target']
```

```
from sklearn.preprocessing import LabelEncoder  
le=LabelEncoder()  
train['sex']=le.fit_transform(train['sex'])  
train['marital.status']=le.fit_transform(train['marital.status'])  
train['education']=le.fit_transform(train['education'])  
train['relationship']=le.fit_transform(train['relationship'])  
train['race']=le.fit_transform(train['race'])  
train['occupation']=le.fit_transform(train['occupation'])  
train['workclass']=le.fit_transform(train['workclass'])  
train['native.country']=le.fit_transform(train['native.country'])  
train.head()
```

```
dd= pd.read_csv('test.csv')
```

```
dd.head()  
dd.shape
```

```
dd.info()  
dd.describe()
```

```
dd.isnull().sum()
```

```
test=dt.dropna()  
test.isnull().sum()
```

```
#encoding
test.columns = ['age', 'workclass', 'fnlwgt', 'education', 'education.num', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'captial.gain', 'capital.loss', 'hours.per.week', 'native.country', 'target']
```

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
test['sex']=le.fit_transform(test['sex'])
test['marital.status']=le.fit_transform(test['marital.status'])
test['education']=le.fit_transform(test['education'])
test['relationship']=le.fit_transform(test['relationship'])
test['race']=le.fit_transform(test['race'])
test['occupation']=le.fit_transform(test['occupation'])
test['workclass']=le.fit_transform(test['workclass'])
test['native.country']=le.fit_transform(test['native.country'])

test.head()
```

```
from sklearn.model_selection import train_test_split
X = train[['age', 'workclass', 'fnlwgt', 'education', 'education.num', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'captial.gain', 'capital.loss', 'hours.per.week', 'native.country']]
y = train['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33)
```

```
#logestic regression
from sklearn.linear_model import LogisticRegression
logmodel = LogisticRegression()
logmodel.fit(X_train,y_train)
valid_predictions = logmodel.predict(X_test)
from sklearn.metrics import classification_report
print(classification_report(y_test,valid_predictions))
```

```
#svm
from sklearn.svm import SVC
svc_model = SVC()
svc_model.fit(X_train,y_train)
svm_predictions = svc_model.predict(X_test)
from sklearn.metrics import classification_report,confusion_matrix
```

```
print(confusion_matrix(y_test,svm_predictions))
print(classification_report(y_test,svm_predictions))
```

#random forest

```
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier(n_estimators=600)
rfc.fit(X_train,y_train)
rf_predictions = rfc.predict(X_test)
print(classification_report(y_test,rf_predictions))
```

#testing

```
rfc = RandomForestClassifier(n_estimators=600)
rfc.fit(X,y)
final_predictions = rfc.predict(test.drop(['target'],axis=1))
print(confusion_matrix(test['target'], final_predictions))
```

OUTPUT:

INFO

```
[ ] dt.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                  32561 non-null  int64
1   workclass            30725 non-null  object
2   fnlwgt               32561 non-null  int64
3   education            32561 non-null  object
4   education.num        32561 non-null  int64
5   marital.status       32561 non-null  object
6   occupation           30718 non-null  object
7   relationship         32561 non-null  object
8   race                 32561 non-null  object
9   sex                  32561 non-null  object
10  capital.gain         32561 non-null  int64
11  capital.loss         32561 non-null  int64
12  hours.per.week       32561 non-null  int64
13  native.country       31978 non-null  object
14  target               32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

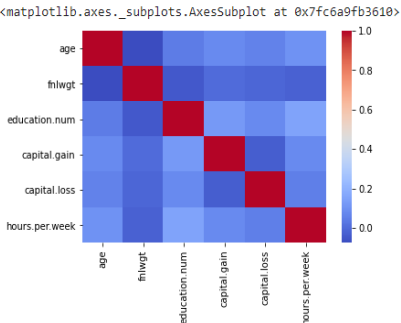
DESCRIBE

```
[ ] dt.describe()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

HEATMAP

```
[ ] sns.heatmap(train.corr(),cmap='coolwarm')
```



ENCODED TRAIN VALUE

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain	capital.loss	hours.per.week	native.country
0	39	5	77516	9	13	4	0	1	4	1	2174	0	40	
1	50	4	83311	9	13	2	3	0	4	1	0	0	13	
2	38	2	215646	11	9	0	5	1	4	1	0	0	40	
3	53	2	234721	1	7	2	5	0	2	1	0	0	40	

INFO

```
dd.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16281 entries, 0 to 16280
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                  16281 non-null  int64
1   workclass            15318 non-null  object
2   fnlwgt               16281 non-null  int64
3   education            16281 non-null  object
4   education.num        16281 non-null  int64
5   marital.status       16281 non-null  object
6   occupation            15315 non-null  object
7   relationship         16281 non-null  object
8   race                 16281 non-null  object
9   sex                  16281 non-null  object
10  capital.gain         16281 non-null  int64
11  capital.loss         16281 non-null  int64
12  hours.per.week       16281 non-null  int64
13  native.country       16007 non-null  object
14  target               16281 non-null  object
dtypes: int64(6), object(9)
memory usage: 1.9+ MB
```

DESCRIBE

dd.describe()

	age	fnlwtg	education.num	capital.gain	capital.loss	hours.per.week
count	16281.000000	1.628100e+04	16281.000000	16281.000000	16281.000000	16281.000000
mean	38.767459	1.894357e+05	10.072907	1081.905104	87.899269	40.392236
std	13.849187	1.057149e+05	2.567545	7583.935968	403.105286	12.479332
min	17.000000	1.349200e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.167360e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.778310e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.383840e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.490400e+06	16.000000	99999.000000	3770.000000	99.000000

ENCODED TEST VALUE

	age	workclass	fnlwtg	education	education.num	marital.status	occupation	relationship	race	sex	captial.gain	capital.loss	hours.per.week	native.cou
0	39	5	77516	9	13	4	0	1	4	1	2174	0	40	
1	50	4	83311	9	13	2	3	0	4	1	0	0	13	
2	38	2	215646	11	9	0	5	1	4	1	0	0	40	
3	53	2	234721	1	7	2	5	0	2	1	0	0	40	

LOGISTIC REGRESSION

print(classification_report(y_test,valid_predictions))

	precision	recall	f1-score	support
<=50K	0.88	0.95	0.87	7479
>50K	0.64	0.29	0.48	2475
accuracy			0.78	9954
macro avg	0.72	0.62	0.63	9954
weighted avg	0.76	0.78	0.75	9954

SVM

[] print(classification_report(y_test,svm_predictions))

	precision	recall	f1-score	support
<=50K	0.78	1.00	0.88	7479
>50K	0.98	0.15	0.26	2475
accuracy			0.79	9954
macro avg	0.88	0.57	0.57	9954
weighted avg	0.83	0.79	0.72	9954

RANDOM FOREST

print(classification_report(y_test,rf_predictions))

	precision	recall	f1-score	support
<=50K	0.89	0.92	0.91	7479
>50K	0.74	0.64	0.69	2475
accuracy			0.85	9954
macro avg	0.81	0.78	0.80	9954
weighted avg	0.85	0.85	0.85	9954

CONCLUSION:

Comparing the accuracy of the different models, Random forest is the best. It gives accuracy of about 85% than compared to other to model.so random forest is better suited for this type of dataset.the salary of the person >50 or not are also shown by each type of model.